

Shibboleth authentication for non-web application

ANDREA BIANCINI, INFN - Sezione Milano Bicocca

andrea.biancini@mib.infn.it

FABIO FARINA, Consortium GARR

fabio.farina@garr.it

SIMON VOCELLA, Consortium GARR

simon.vocella@garr.it

Abstract Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum. Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum. Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum. Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Categories and Subject Descriptors: C.2.2 [**Computer-Communication Networks**]: Network Protocols

General Terms: Shibboleth, PAM, JAAS

Additional Key Words and Phrases: Shibboleth, SAML, Basic authentication, PAM, JAAS, Python

ACM Reference Format:

Biancini A., Farina F., Vocella S. 2012. Shibboleth authentication for non-web application. ACM Trans. Embedd. Comput. Syst. 9, 4, Article 39 (March 2010), 11 pages.

DOI = 10.1145/0000000.0000000 <http://doi.acm.org/10.1145/0000000.0000000>

1. INTRODUCTION

User authentication and user authorization are main tasks needed to guarantee security in IT environments. One important requirement of such security solutions, is that they must be as transparent as possible. Current user applications, usually sit on different systems and environment. This requires a solution able to permit that each system is not required to create and manage its own credential system. For this reason single sign-on (SSO, as described in [Shaer 1995]) systems have been designed and implemented. These systems have also a strong impact on users who don't need to perform multiple logins on different systems, but could be automatically identified by all the application they are using.

One of the most widespread SSO software package is Shibboleth [Morgan et al. 2004], which implements authentication and authorization over network resources with a federated approach. Shibboleth, designed and implemented by Internet2, provides an effective solution for secure multi-organizational access to web resources. It implements widely used federated identity standards, principally OASIS' Security Assertion Markup Language (SAML [Cantor et al. 2005]), to provide a federated single sign-on and attribute exchange framework. Shibboleth also provides extended privacy functionality allowing the browser user and their home site to control the attributes released to each application.

Using Shibboleth-enabled access simplifies the management of identity and permissions for organizations supporting either users or applications. Shibboleth is developed in an open and participatory environment and is freely available. This open and collaborative nature has allowed Shibboleth to become the heart of many solutions for single sign-on and Authentication & Authorization Infrastructures (AAI) in several identity federations.

The architecture of Shibboleth, described in [Erdos and Cantor 2005], identifies two main components: Identity Providers (IdPs) which authenticate the users and supply user authorization attributes; Service Providers (SPs) which consume user attributes and provide access to the secure contents.

The communication exchanges between the user, the IdP and the SP are all based on SAML assertions and transit over HTTP channels. This implementation architecture permits Shibboleth to work seamlessly inside user web-browsers. However, this approach is strongly web-centric, and for this reason does not adapt well to provide SSO authentication to non web-based application.

In this article a solution is presented to permit Shibboleth authentication for non web-based applications. The traditional login scheme used by Shibboleth to authenticate users is implemented inside a web-browser. The user, when trying to log to a web resource, is redirected to the IdP to perform authentication and obtain authorization assertions via his session metadata. The use case implemented within this article solves the problem of authenticating a user and retrieving his metadata information from Shibboleth without using a web-browser. This solution could be integrated in traditional client-based application (for instance application written in Java or Python) to obtain a secure way to manage user logins.

This article will describe the architecture and the technical implementation of this authentication solution. As an example of possible uses of this work, different software implementations will be described showing how to integrate the Shibboleth

AAI mechanisms inside several client applications (Linux platform, Java JAAS, Python).

The rest of this paper is organized as follows: the first section contains a quick overview of related works; the second section will present the architecture of the solution presented with this article; the third section will present some details about the technical implementation in the different software platform; the last section presents conclusions and future works.

2. RELATED WORKS

To solve this problem of AAI integration between Shibboleth and non-web application the more general and wide approach is the Moonshot project ([Howlett and Hartman 2005]). This project aims explicitly to develop a single unifying technology for extending the benefits of federated identity to a broad range of non-Web services. The Moonshot project proposes a solution to the AAI problem very complete and articulated.

It leverages different technologies such as Kerberos (a computer network authentication protocol which allows nodes communicating over a non-secure network to prove their identity to one another in a secure manner), the GSS library (Generic Security Services libraries for programs to access security services) and a Radius server (Remote Authentication Dial In User Service, a networking protocol that provides centralized Authentication, Authorization, and Accounting management).

This solution is for sure very complete and trustworthy. It implies however a quite complex architecture and requires a lot of changes on client side, SPs and IdPs. This suggests a difficult introduction of this approach into real existing federations.

The need to leverage Shibboleth identity federations in non-web application has been particularly felt in Grid Computing ([Kesselman and Foster 1998]) services. In the field of research, in fact, in the last years strong investments have been done to build a global e-Infrastructure leveraging Grid Computing technologies. This infrastructure supports the execution of scientific computation and the storage of relevant research data.

The grid, since its inception, has introduced specific solutions (inspired to the original article [Foster et al. 1998]) to security problems. The new Shibboleth SSO for web-application somehow lies next to the Grid AAI, thus the idea to integrate the two security systems is natural and potentially very beneficial.

Due to the complexity of the Moonshot project, which would add to the intrinsic complexity of grid security schemas, different solutions have been adopted in this field.

In grid environments, this problem has been approached with the goal to try and integrate the existing grid security infrastructures with Shibboleth mechanisms. These solutions, described in both articles [Wang et al. 2009] and [Jensen et al. 2007], permit the users to transparently move between Shibboleth application and Grid environments using one single set of credentials. They work with specific software middlewares able to log in the user and then manage both the Shibboleth authentication assertions and the grid certificates used to access the different environments where applications sit.

Solutions like those, are very interesting in that they really ease the user experience and provide effective results. However it must be noted that are well fit for the problem of integrating Shibboleth with grid security and can hardly be extended to other fields of application. They are not able to provide a solution for the AAI needs of non-web application not running on a grid environment.

The approaches described show opposite characteristics. On one side a very complex and articulated mechanism has been proposed, on the other side very specific and poorly adaptable solutions have been proposed to solve pragmatic problems.

In this paper a different approach will be presented. The proposed approach would try and address general problems of authentication, but will try to do it with few interventions on the requirements of SPs and IdPs.

3. ARCHITECTURE

The solution proposed leverages the standard authentication and authorization mechanisms implemented in Shibboleth. It executes a login to the IdP using HTTP Basic mechanism ([Franks et al. 1999]) over https. The authentication process of Shibboleth works as described in figure 1.

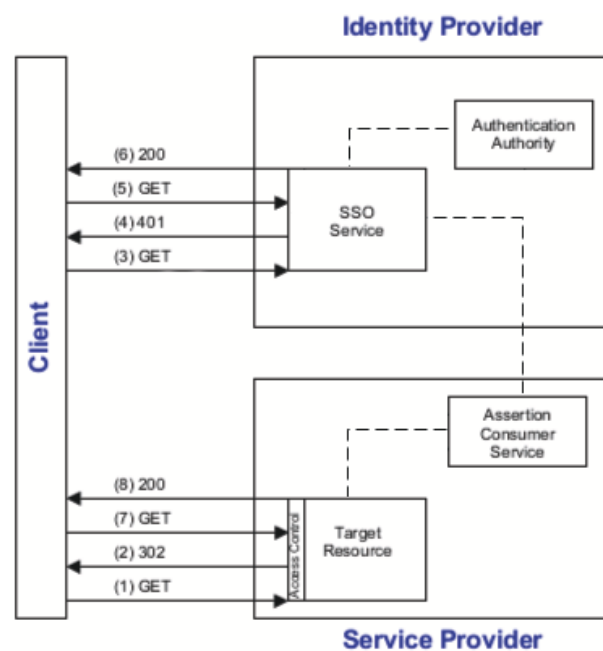


Fig. 1. Shibboleth login procedure

The solution described in this article follows the same schema, but with the following clarifications:

- (1) The original request is directed to a web-page on the SP side that is behind Shibboleth authentication. This page will produce a list of rows key=value containing the metadata information to be passed to user session after login.
- (2) The webserver on the SP answers with a redirect operation to a page on the IdP.
- (3) The client follows the redirect and opens the IdP login page.
- (4) This request obtains an answer “authorization required” asking the client to authenticate via HTTP Basic.
- (5) The client authenticates performing the same request with the proper HTTP header containing username and password for the user. The answer to this request contains a Set-Cookie instruction with which the IdP creates a cookie on the client

ALGORITHM 1: Example of PAM webpage on the S

```

authenticated=true
Shib-Application-ID=default
Shib-Session-ID=_d28ec2af03fd7437ca500c351ab0aec0
Shib-Identity-Provider=http://idp.server.hosntame/idp/shibboleth
Shib-Authentication-Instant=2012-06-15T12:55:45.587Z
Shib-Session-Index=0141d3c23dc19ee047f68e35037828a9db68d7961082566c9b399e512e03705f
eduPersonEntitlement=urn:mace:mib.infn.it:permission:service1:access:user
eduPersonPrincipalName=user@domain.it
eduPersonScopedAffiliation=member@domain.it;student@domain.it
eduPersonTargetedID=http://idp.server.hosntame/idp/shibboleth!https://sp.server...
email=name.surname@domain.it
givenName=Name
sn=Surname
uid=name
Shib-Session-Unique=64656661756c7468747470733a2f2f636c6f75642d6d692d30332e6d696...

```

containing a couple of keys permitting the client to reconnect with the Shibboleth session created.

- (6) The IdP communicates with the SP to create a valid session for the user with all the metadata of the user session.
- (7) The client opens the original URL providing the cookie set by the IdP after the authentication. This request gets the page with the key=value rows used to initialize the user session on the client side.

This mechanism simulates what usually happens in the browser of a user authentication to a Shibboleth web-application. All this interactions are, however, automatic and do not require a user to manage or control the underlying HTTP dialogue to permit the authentication.

In order for this mechanism to work, a web page has been created on a server and placed behind Shibboleth authentication. The user, after having provided valid login information, is shown this page that consists of a list of rows with the format key=value. This list represents the user session and shows all the relevant metadata provided by Shibboleth with the user authentication. An example of the content of this page is as follows:

4. IMPLEMENTATIONS

This article will present different implementations of this AAI solution. In order to test it in different scenarios and a wide spectrum of application, a set of different implementations have been realized covering the more widespread application platforms.

In particular three main implementations have been realized: an integration of these authorization and authentication mechanisms inside Linux PAM and NSS modules; a integration with Java JAAS architecture; a module for python applications. In the following part of this section, these implementations will be described.

4.1. PAM and NSS modules

The standard mechanisms used by modern Linux systems to authenticate and list users are PAM and NSS. These two libraries have different purposes and implement different functionalities.

4.1.1. *Pluggable Authentication Modules (PAM)*. Pluggable Authentication Modules (PAM) are at the core of user authentication in any modern linux distribution. PAM enables the user programs to transparently authenticate users, regardless of how security information is stored. As described in the Linux-PAM System Administrator's Guide ([Morgan and Kukuk 1996]): "It is the purpose of the Linux-PAM project to separate the development of privilege granting software from the development of secure and appropriate authentication schemes. This is accomplished by providing a library of functions that an application may use to request that a user be authenticated".

For this reason, in order to guarantee a SSO mechanism on linux machines it seemed straightforward to implement a module for PAM able to authenticate and retrieve user sessions from Shibboleth via the HTTP Basic authentication mechanism. In order to authenticate using HTTP Basic, the general schema presented in section 2 is implemented using the libcurl library ([Stenberg 1996]).

Libcurl is used to manage all the HTTP dialogues and is able to manage cookies and perform the HTTP basic authentication, when required by the IdP.

A PAM module is a shared library with a public interface, specified by PAM headers, that links with the user programs to perform actions of authentication, accounting and to retrieve user session values. The methods that needs to be implemented are the following:

- `pam_sm_authenticate`: this function performs user authentication. This function has to request to the user his user ID and password. To perform these operation the PAM library uses an application-defined callback to allow a direct communication between a loaded module and the user via the application. This function, retrieved, the needed information, starts the login procedure with the SP and IdP and provides an answer whether the user has been able to provide a valid authentication or not.
- `pam_sm_setcred`: this function retrieves information to create a session for the user. Generally, an authentication module may have access to more information about a user than their authentication token. This function is used to make such information available to the application. The session values are obtained from the web page used to perform HTTP Basic authentication. As described in section 2, this page contains in fact a list of key=value rows retrieving information from Shibboleth metadata and describing the logged user session.
- `pam_sm_acct_mgmt`: this function performs the task of establishing whether the user is permitted to gain access at this time. This function gets called when the user has already been validated by an authentication module. This function checks for the returned page body of the HTTP target resources behind Shibboleth authentication to verify that the `authenticate` variable is specified with value `true`.
- `pam_sm_open_session`: this function is called to commence a session.
- `pam_sm_close_session`: this function is called to terminate a session.
- `pam_sm_chauthtok`: this function is used to (re-)set the authentication token of the user. This function is not implemented in the module realized for Shibboleth. In fact the PAM module realized is not able to modify user information on the IdP side but it is used only to check user session and perform AAI.

This PAM module developed, can then be used to authenticate users in a common PAM chain. For example, modifying the file `/etc/pam.d/login` and modifying with the config file in algorithm 2, it is possible to permit login to the linux machine using the Shibboleth credentials.

The `pam_http.so` module can be instantiated with some parameters. In particular, it is possible to specify the following parameters:

ALGORITHM 2: PAM configuration to use the pam_http module.

```

auth required pam_http.so url=https://servername.com/pam sess_username=username
account required pam_http.so
password required pam_permit.so
session required pam_http.so

```

- url: specifies the URL of the page behind Shibboleth authentication on the SP that produces as output the data to be put in user session after login;
- sess_username: specifies the field in the user session containing the username to be passed to the system for the logged in user; this name, in the standard linux process, can differ from the one used for login;
- sslcheck (optional): specifies whether the libcurl library has to verify the SSL certificate of the server when connecting via HTTPS;
- cafile (optional): specifies the name of a file, containing a CA public key, used by libcurl to verify the SSL certificate of the server when connecting via HTTPS.

4.1.2. *Name Service Switch (NSS)*. To permit a linux system to user external directories for users and user groups, it is necessary to implement another module for a different library: the Name Service Switch (NSS, [Foundation 1987]). This library permits to define services for accessing different databases containing various information. In respect of the activities described in this article, the relevant databases are: passwd, the user database; and group a database with all user groups.

The service implemented to integrate NSS with Shibboleth is a new servlet that has been deployed on the IdP. This servlet connects to the underlying authorization database (usually an LDAP database) and retrieves information about all users and groups that are recognized by that IdP.

The NSS module implemented, then, had to implement the following functions:

- `_nss_shib_getpwnam_r`: this function is called to retrieve user information starting from his user-name;
- `_nss_shib_getpwuid_r`: this function is called to retrieve user information starting from his uid;
- `_nss_shib_getpwent_r`: this function is called to retrieve all the available users;
- `_nss_shib_getgrnam_r`: this function is called to retrieve a user group information starting from its name;
- `_nss_shib_getgrgid_r`: this function is called to retrieve a user group information starting from its id (gid);
- `_nss_shib_getgrent_r`: this function is called to retrieve all the available user groups.

This NSS module can be activated by modifying the `/etc/nsswitch.conf` file as shown in algorithm 3. The NSS module reads some configuration from the configuration file `/etc/libnss.conf`. In particular, it is possible to specify the following parameters:

- url: specifies the URL of the servlet created on the IdP and able to list users and groups of users of a specific identity provider;
- sslcheck (optional): as for the PAM module, specifies whether the libcurl library has to verify the SSL certificate of the server when connecting via HTTPS;

ALGORITHM 3: NSS configuration to use the shib module.

```
passwd: files shib
shadow: files
group: files shib
```

```
hosts: files dns
```

- `cafile` (optional): as for the PAM module, specifies the name of a file, containing a CA public key, used by libcurl to verify the SSL certificate of the server when connecting via HTTPS;
- `username` (optional): if the servlet is behind an authentication method, specifies the username to be used for login;
- `password` (optional): if the servlet is behind an authentication method, specifies the password to be used for login;
- `cookie_num`: specifies the number of cookies to be passed to the servlet; if no cookie has to be passed, 0 must be specified here;
- `cookie_#_name`: specifies for each cookie the name of the cookie to be passed; values with `${name}` are considered variables to be searched in the environment;
- `cookie_#_value`: specifies for each cookie the value of the cookie to be passed; values with `${name}` are considered variables to be searched in the environment.

With this two modules the linux machine recognizes Shibboleth users in the same way as local users. So a Shibboleth user can login to the linux machine with his credentials, can have access rights on files, use ACLs, and so on...

After the login the user finds in its login environment all the Shibboleth session values. All metadata downloaded from Shibboleth is in fact placed in the user session and can be used to perform SSO with other non web-based application.

4.2. JAAS module

Java, with its Enterprise Edition, has become a very well known and ubiquitous platform for developing software applications. Starting from Java 1.4 a common service has been introduced for AAI: the Java Authentication and Authorization Service (JAAS, [Java 2002]). JAAS can be used for two purposes:

- for authentication of users, to reliably and securely determine who is currently executing Java code, regardless of whether the code is running as an application, an applet, a bean, or a servlet; and
- for authorization of users to ensure they have the access control rights (permissions) required to do the actions performed.

JAAS implements a Java version of the standard Pluggable Authentication Module (PAM) framework. For this reason, and for its widespread utilization in real applications, JAAS seemed a proper candidate to integrate a module to authenticate users on Shibboleth. Via this module it is now possible to integrate Shibboleth authentication, and metadata management, inside Java application not based on web-technologies.

The general architecture of this JAAS module is the same of the PAM module described in this article in section 4.1. A JAAS module is implemented by a class extending the `javax.security.auth.spi.LoginModule` class and implementing the following methods:

ALGORITHM 4: JAAS configuration file to use the JAASShibbolethLoginModule login module.

```
Shibboleth {
  it.infn.mib.shibboleth.jaas.JAASShibbolethLoginModule required
  url="https://servername.com/pam"
  sslcheck="false"
  sess_username="username"
  truststore=""
  truststore_password=""
  debug="false";
};
```

ALGORITHM 5: Example of a Java login with JAASShibbolethLoginModule.

```
LoginContext lc = new LoginContext("Shibboleth", new MyCallbackHandler());
lc.login();
System.out.println("User logged in successfully.");

Map<String, String> session = null; // will contain the Shibboleth user session
for (Principal curPrincipal : lc.getSubject().getPrincipals()) {
  if (curPrincipal instanceof ShibbolethPrincipal) {
    loggedUser = ((ShibbolethPrincipal) curPrincipal).getName();
    session = ((ShibbolethPrincipal) curPrincipal).getSession();
  }
}
```

- initialize: this method is used to initialize the login module specifying: the subject to be authenticated, a callback handler to deal the communications with the end user (in a very similar way to PAM conversation mechanisms);
- login: this method authenticates the user by checking his user name and password;
- logout: this method logs out the user;
- abort: this method is called if the login process terminates for some reason internal to JAAS configuration and cleans up any state that was originally saved;
- commit: this method is called if the login process concludes positively; it associates a set of principals to the login context; these principals can contain relevant information about the user session.

Within the JAAS module for Shibboleth a specific principal has been developed (the `ShibbolethPrincipal`). This principal contains a `Map<String, String>` object that represents the user session containing all Shibboleth metadata information.

This JAAS module can be configured with a configuration file as the one presented in algorithm 4. The login module accepts the same parameter of the PAM module described in section 4.1 but, instead of having `cafile`, has a mechanism to authenticate server SSL certificates that leverages the standard truststores of Java defined by specifying the truststore file name and the password to access it.

4.2.1. Example of login and credential usage in a Java application. In order to integrate such authentication mechanism into a real Java application, very few lines of code have to be written. The algorithm 5 shows how it is possible to use the Shibboleth login context defined in the JAAS config file to authenticate the user and then how it is possible to cycle through all the principals returned by the login process to retrieve the user session.

After the login, it is possible to use the credentials obtained to connect to other services over the network. For instance, it could be possible to call a webservice behind

the same SP connecting to the Shibboleth session created during login.

In order to connect to an existing session, it is sufficient to pass to the web-service a cookie as follows:

- name: the name of the cookies is ‘_shibsession_’ followed by a unique string identifying the session. This string can be obtained from the user session via the variable called Shib-Session-Unique.
- value: the value of this cookie is the shibboleth session id for the user’s session. This value can be obtained from the user session via the variable called Shib-Session-ID.

The two variables used to create the cookie entry to be passed to HTTP webservice call, can be retrieved from the JAAS Shibboleth principal or from the user shell environment, if the user logged the linux machine using the PAM and NSS modules described within this article.

4.3. Python

5. CONCLUSION AND FUTURE WORKS

REFERENCES

- CANTOR, S., KEMP, J., PHILPOTT, R., AND MALER, E. 2005. Assertions and protocols for the OASIS security assertion markup language (SAML) v2.0. Tech. rep. 03.
- ERDOS, M. AND CANTOR, S. 2005. The Shibboleth architecture. <http://shibboleth.internet2.edu/>.
- FOSTER, I., KESSELMAN, C., TSUDIK, G., AND TUECKE, S. 1998. A security architecture for computational grids. In *Proceedings of the 5th ACM conference on Computer and communications security*. CCS '98. ACM, 83–92.
- FOUNDATION, T. F. S. 1987. Name service switch. <http://www.gnu.org/>.
- FRANKS, J., HALLAM-BAKER, P., HOSTETLER, J., LAWRENCE, S., LEACH, P., LUOTONEN, A., AND STEWART, L. 1999. Http authentication: Basic and digest access authentication. RFC 2617.
- HOWLETT, J. AND HARTMAN, S. 2005. Project moonshot. Tech. rep. 07.
- JAVA. 2002. Java authentication and authorization service. <http://java.sun.com/products/jaas/>.
- JENSEN, J., WALLOM, D., SPENCE, D., TANG, K., MEREDITH, D., AND TRETHEFEN, A. 2007. Shibgrid, a shibboleth based access method for the national grid service. In *UK e-Science All Hands Meeting*.
- KESSELMAN, C. AND FOSTER, I. 1998. *The Grid: Blueprint for a New Computing Infrastructure*. Morgan Kaufmann Publishers.
- MORGAN, A. G. AND KUKUK, T. 1996. The linux-pam system administrators' guide. <http://kernel.org/>.
- MORGAN, R. L., CANTOR, S., CARMODY, S., HOEHN, W., AND KLINGENSTEIN, K. 2004. Federated security: The shibboleth approach. *EDUCAUSE Quarterly* 27, 4, 12–17.
- SHAER, C. 1995. Single sign-on. *Network Security 1995*, 8, 11–15.
- STENBERG, D. 1996. curl and libcurl. <http://curl.haxx.se/>.
- WANG, X. D., JONES, M., JENSEN, J., RICHARDS, A., WALLOM, D., MA, T., FRANK, R., SPENCE, D., YOUNG, S., DEVEREUX, C., AND GEDDES, N. 2009. Shibboleth access for resources on the national grid service (sarongs). In *Proceedings of the 2009 Fifth International Conference on Information Assurance and Security - Volume 02*. IAS '09. IEEE Computer Society, 338–341.

Received September 2012; revised September 2012; accepted September 2012